


What explains the results



Best arm per model (1 = best) agreement between models $r = +0.11$

	distributional	scalar	scalar_cvar5	placebo	placebo_shuffled
deepseek_v4_pro	3	1	2	4	5
gemini_2_5_flash	5	2	1	3	4
glm_5_2	4	5	1	2	3
gpt_5_6_luna	5	3	2	4	1
haiku_4_5	2	5	4	1	3
kimi_k3	2	3	5	1	4
nemotron_3_super	4	5	2	1	3
opus_5	5	3	1	2	4
qwen3_5_9b	5	3	1	2	4
qwen3_6_27b	2	4	3	1	5
sonnet_5	2	1	4	3	5